DAY 32 | PRE-READ

Running AI Like a Product

A light read before Day 32. About 15 minutes, no code required.

Prepared by **Ahmed** | linkedin.com/in/ahmed-al-imaad

Why read this?

Tomorrow you learn how teams keep AI systems accurate, affordable and fast over months, not just on launch day. There is also a 1.5 hour assessment.

1. Teaching new habits

A new employee can follow written instructions (prompting) or look things up in a manual (RAG). But if you want them to always write reports in your house style without being reminded, you train them with many good examples. That is **fine-tuning**: teaching a model new habits from examples. It is powerful, but it does not teach fresh facts; documents and RAG do that better.

2. Knowing exactly what you shipped

A pharmacy records the batch number of every medicine, so if there is a problem it knows exactly which batch to recall. AI teams do the same: they record the exact prompt, model version, search settings and test data behind every release. This is **versioning**.

3. Every word has a price

AI services usually charge by the amount of text processed, counted in tokens. Long prompts, long answers and large models all cost more and take longer. Teams cut cost by keeping prompts lean, reusing answers to repeated questions, and sending easy questions to smaller, cheaper models.

Technique	Everyday comparison
Caching	Keeping a printed FAQ instead of asking the expert every time
Routing	A receptionist handles simple questions; specialists get the hard ones
Trimming context	Sending one relevant page, not the whole binder

Key idea: measure cost and speed for every request. You cannot manage what you do not measure.

4. Think about these before class

- Which questions does your team answer again and again? Could they be cached?
- What would you need to record to explain an AI answer given six months ago?

- Which tasks in your work follow a strict format that a fine-tuned model could learn?

Kiroku (記録): a record, keeping records