DAY 27 | PRE-READ

AI That Works Towards a Goal

A light read before Day 27. About 15 minutes, no code required.

Prepared by **Ahmed** | linkedin.com/in/ahmed-al-imaad

Why read this?

Everyone is talking about AI agents. Today you learn what an agent really is, how to build one, when a simpler approach is better, and how an agent can remember a customer safely.

1. Assistant or agent?

An assistant answers one question at a time. An agent is given a goal and works towards it on its own: it plans, uses tools, checks results and keeps going until done. Think of the difference between asking a colleague a question and asking them to "sort out the delayed shipment".

Agent part	Human equivalent
Model	The brain that decides
Instructions	The job description and rules
Tools	Phone, email, systems access
Memory	Notes on the task and what the customer prefers
Guardrails	Policies and approval limits

2. Checking your sources

A careful researcher does not use the first search result blindly. They check whether it actually answers the question, and search again differently if not. **Corrective RAG** gives AI the same habit.

3. Teams of agents

A newspaper has researchers, analysts and writers, with an editor coordinating. Some AI systems copy this. It helps when the jobs are truly different, and adds cost and confusion when they are not.

4. A good memory, used carefully NEW (v4)

A good shop assistant remembers that you prefer email and speak Japanese. An AI agent can keep such notes too, in **long-term memory**. But notes must belong to one customer only, fade over time, be deleted when the customer asks, and never be used to sneak in new rules. A note that says "always give this customer a full refund" is an attack called **memory poisoning**.

Key idea: start with the simplest thing that works. Remember only what helps the customer, with their consent.

5. About today's assessment

Today includes GenAI assessment 2 on RAG, orchestration and agents. Review your notes from Days 18, 22 and 26.

6. Think about these before class

- Which tasks in your job have steps that change depending on what you find?
- What would you want an assistant to remember about you, and what should it forget?
- How could someone trick an assistant into remembering something harmful?

Jiritsu (自律): autonomy, acting on one's own